

```
In [22]: !pip install imblearn
```

```
Defaulting to user installation because normal site-packages is not writeable
Collecting imblearn
  Downloading imblearn-0.0-py2.py3-none-any.whl (1.9 kB)
Collecting imbalanced-learn
  Downloading imbalanced_learn-0.10.1-py3-none-any.whl (226 kB)
----- 226.0/226.0 kB 655.8 kB/s eta 0:00:00
Requirement already satisfied: joblib>=1.1.1 in c:\users\jkuma\appdata\roaming\python\python39\site-packages (from imbalanced-learn->imblearn) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\programdata\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (2.2.0)
Requirement already satisfied: scipy>=1.3.2 in c:\programdata\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.9.1)
Requirement already satisfied: numpy>=1.17.3 in c:\programdata\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.21.5)
Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\jkuma\appdata\roaming\python\python39\site-packages (from imbalanced-learn->imblearn) (1.2.2)
Installing collected packages: imbalanced-learn, imblearn
Successfully installed imbalanced-learn-0.10.1 imblearn-0.0
```

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.linear_model import LogisticRegressionCV
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from imblearn.over_sampling import SMOTE
import plotly.express as px
from sklearn.preprocessing import LabelEncoder
from imblearn.under_sampling import NearMiss
```

```
In [2]: df=pd.read_csv(r"C:\Users\Jkuma\Downloads\heart_2020_cleaned.csv\keyfactorofheartdisease
```

```
In [3]: df
```

Out[3]:	HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth	DiffWalking	Sex
0	No	16.60	Yes	No	No	3.0	30.0	No	Female
1	No	20.34	No	No	Yes	0.0	0.0	No	Female
2	No	26.58	Yes	No	No	20.0	30.0	No	Male
3	No	24.21	No	No	No	0.0	0.0	No	Female
4	No	23.71	No	No	No	28.0	0.0	Yes	Female
...
319790	Yes	27.41	Yes	No	No	7.0	0.0	Yes	Male
319791	No	29.84	Yes	No	No	0.0	0.0	No	Male
319792	No	24.24	No	No	No	0.0	0.0	No	Female
319793	No	32.81	No	No	No	0.0	0.0	No	Female
319794	No	46.56	No	No	No	0.0	0.0	No	Female

319795 rows × 18 columns

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 319795 entries, 0 to 319794
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   HeartDisease          319795 non-null object
1   BMI                   319795 non-null float64
2   Smoking               319795 non-null object
3   AlcoholDrinking       319795 non-null object
4   Stroke                319795 non-null object
5   PhysicalHealth        319795 non-null float64
6   MentalHealth          319795 non-null float64
7   DiffWalking           319795 non-null object
8   Sex                   319795 non-null object
9   AgeCategory           319795 non-null object
10  Race                  319795 non-null object
11  Diabetic              319795 non-null object
12  PhysicalActivity       319795 non-null object
13  GenHealth             319795 non-null object
14  SleepTime             319795 non-null float64
15  Asthma                319795 non-null object
16  KidneyDisease         319795 non-null object
17  SkinCancer            319795 non-null object
dtypes: float64(4), object(14)
memory usage: 43.9+ MB
```

In [5]: `df.columns`

Out[5]: Index(['HeartDisease', 'BMI', 'Smoking', 'AlcoholDrinking', 'Stroke', 'PhysicalHealth', 'MentalHealth', 'DiffWalking', 'Sex', 'AgeCategory', 'Race', 'Diabetic', 'PhysicalActivity', 'GenHealth', 'SleepTime', 'Asthma', 'KidneyDisease', 'SkinCancer'], dtype='object')

In [6]: `df.shape`

Out[6]: (319795, 18)

In []:

In []:

Before training process for model,select only those columns that would significantly impact the likelihood of heart disease

In [7]: `new_df=df[['HeartDisease', 'BMI', 'Smoking', 'AlcoholDrinking', 'Stroke','PhysicalHealth`

In [8]: `new_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 319795 entries, 0 to 319794
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   HeartDisease          319795 non-null object
1   BMI                   319795 non-null float64
2   Smoking               319795 non-null object
3   AlcoholDrinking       319795 non-null object
4   Stroke                319795 non-null object
5   PhysicalHealth        319795 non-null float64
6   MentalHealth          319795 non-null float64
7   DiffWalking           319795 non-null object
8   Sex                   319795 non-null object
9   AgeCategory           319795 non-null object
10  Diabetic              319795 non-null object
11  PhysicalActivity       319795 non-null object
12  GenHealth              319795 non-null object
13  SleepTime             319795 non-null float64
14  Asthma                319795 non-null object
15  KidneyDisease          319795 non-null object
dtypes: float64(4), object(12)
memory usage: 39.0+ MB
```

In [9]: `new_df.describe().T`

Out[9]:

	count	mean	std	min	25%	50%	75%	max
BMI	319795.0	28.325399	6.356100	12.02	24.03	27.34	31.42	94.85
PhysicalHealth	319795.0	3.371710	7.950850	0.00	0.00	0.00	2.00	30.00
MentalHealth	319795.0	3.898366	7.955235	0.00	0.00	0.00	3.00	30.00
SleepTime	319795.0	7.097075	1.436007	1.00	6.00	7.00	8.00	24.00

In [10]: `new_df.isnull().sum()`

```
Out[10]: HeartDisease      0
        BMI              0
        Smoking          0
        AlcoholDrinking  0
        Stroke           0
        PhysicalHealth    0
        MentalHealth      0
        DiffWalking      0
        Sex              0
        AgeCategory      0
        Diabetic          0
        PhysicalActivity  0
        GenHealth         0
        SleepTime         0
        Asthma            0
        KidneyDisease     0
        dtype: int64
```

```
In [11]: new_df.nunique()
```

```
Out[11]: HeartDisease      2
        BMI              3604
        Smoking          2
        AlcoholDrinking  2
        Stroke           2
        PhysicalHealth    31
        MentalHealth      31
        DiffWalking      2
        Sex              2
        AgeCategory      13
        Diabetic          4
        PhysicalActivity  2
        GenHealth         5
        SleepTime         24
        Asthma            2
        KidneyDisease     2
        dtype: int64
```

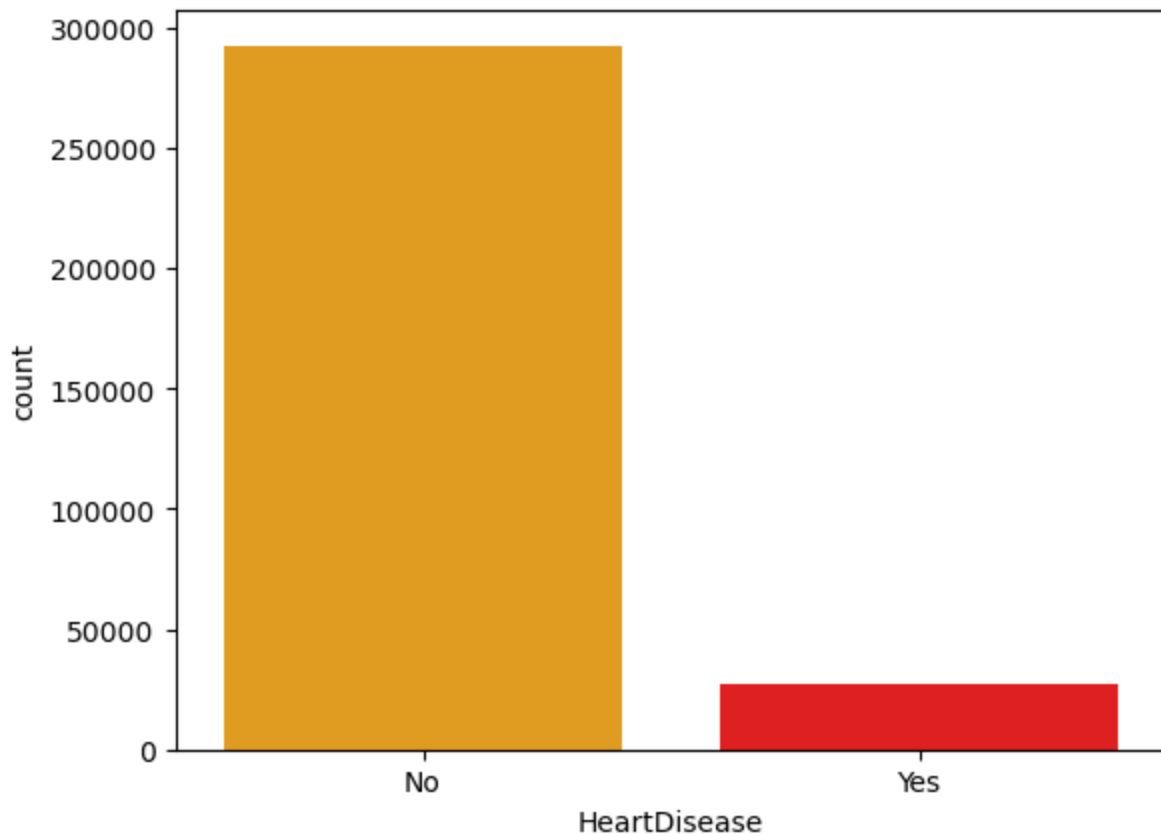
```
In [12]: new_df.dtypes
```

```
Out[12]: HeartDisease      object
        BMI              float64
        Smoking          object
        AlcoholDrinking  object
        Stroke           object
        PhysicalHealth    float64
        MentalHealth      float64
        DiffWalking      object
        Sex              object
        AgeCategory      object
        Diabetic          object
        PhysicalActivity  object
        GenHealth         object
        SleepTime         float64
        Asthma            object
        KidneyDisease     object
        dtype: object
```

Data Visualization

```
In [13]: sns.countplot(x=new_df['HeartDisease'],palette=["Orange", "Red"])
```

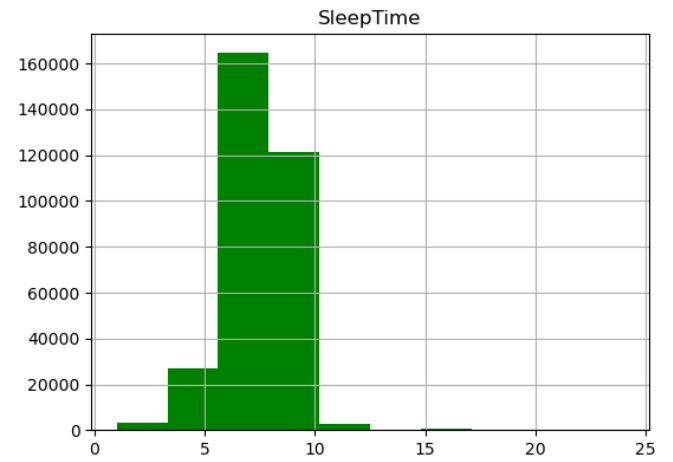

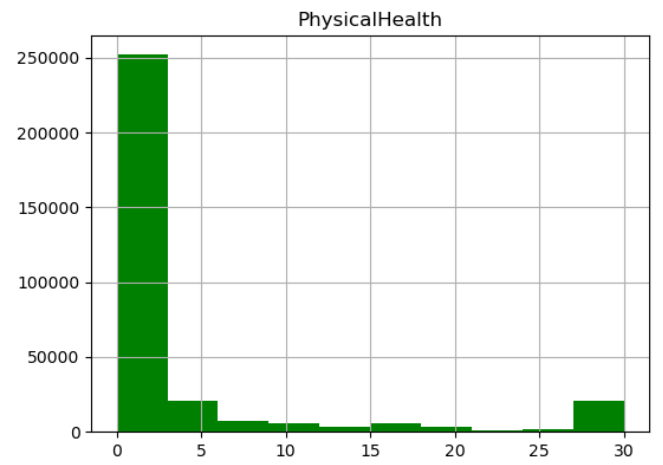
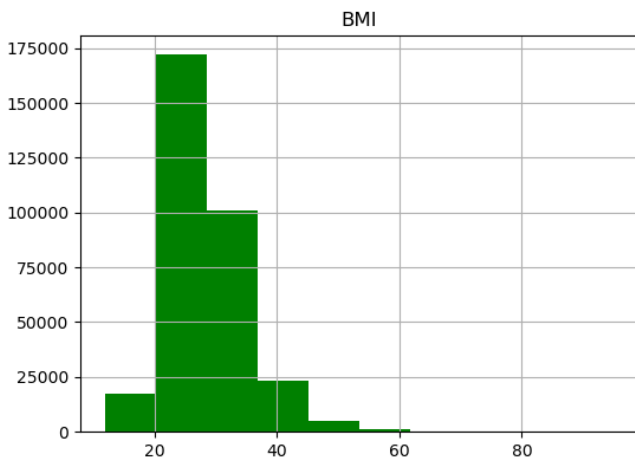
```
Out[13]: <AxesSubplot: xlabel='HeartDisease', ylabel='count'>
```



The graph shows that the amount of records for heart disease are unbalanced, showing an uneven distribution of information between people who have the condition and those who do not. The forecasts for heart disease made by the trained model could be significantly biased as a result. We will employ SMOTE (Synthetic Minority Oversampling Technique) to balance the class distribution. However, in order for SMOTE to work properly, all categorical data must first be converted into binary using dummy variables.

The nearest minority class data elements are taken into account by the SMOTE algorithm, which then generates new combinations based on those entries.

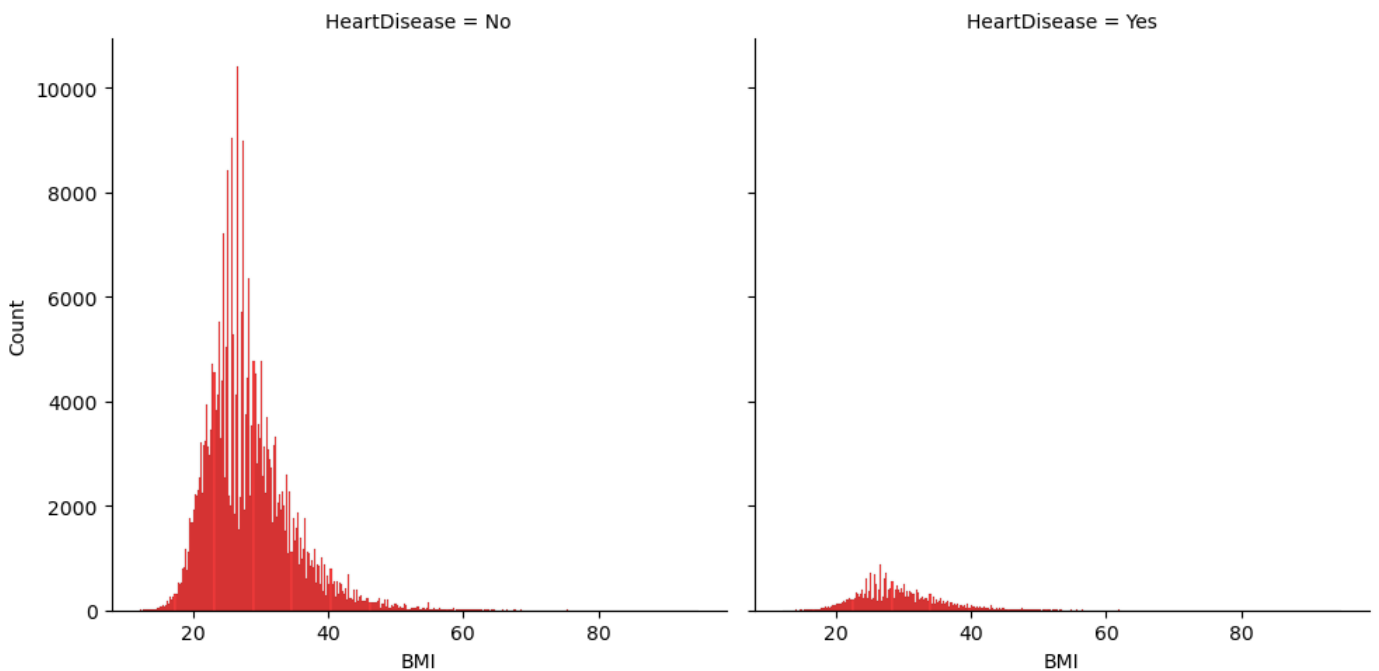
```
In [14]: new_df.hist(bins=10,figsize=(14,10),color="green")  
  
plt.savefig('histogram1.png')
```



In []:


```
In [15]: sns.displot(df,x="BMI",col='HeartDisease',color="red")
```

```
Out[15]: <seaborn.axisgrid.FacetGrid at 0x27404fe7490>
```




```
In [16]: fig = px.pie(new_df,names='Smoking',title='Smoking',color_discrete_sequence=px.colors.se
fig.show()
#
```

Smoking

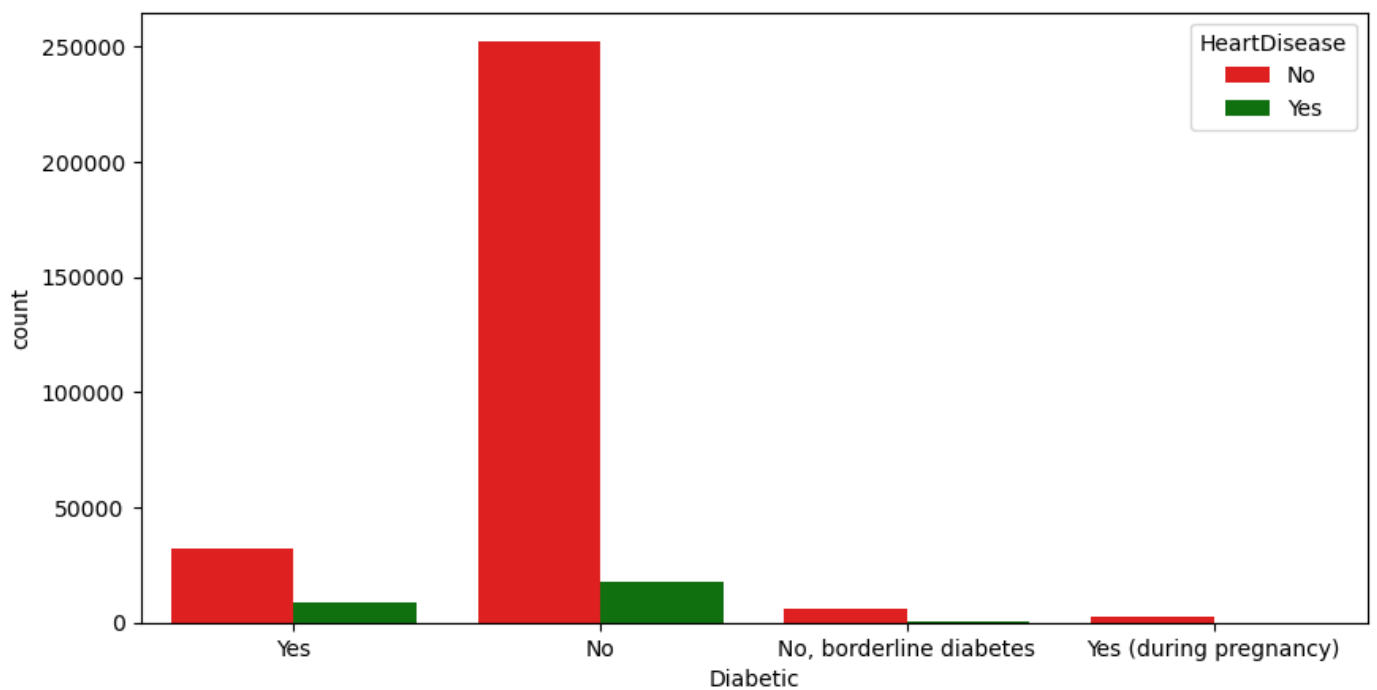


```
In [17]: fig = px.pie(new_df, names='Diabetic', title='Diabetic', color_discrete_sequence=px.colors.  
fig.show()
```

Diabetic



```
In [18]: plt.figure(figsize=(10,5))
sns.countplot(x=new_df['Diabetic'],hue="HeartDisease",data=new_df,palette=[ "Red", "Green"]
plt.show()
```



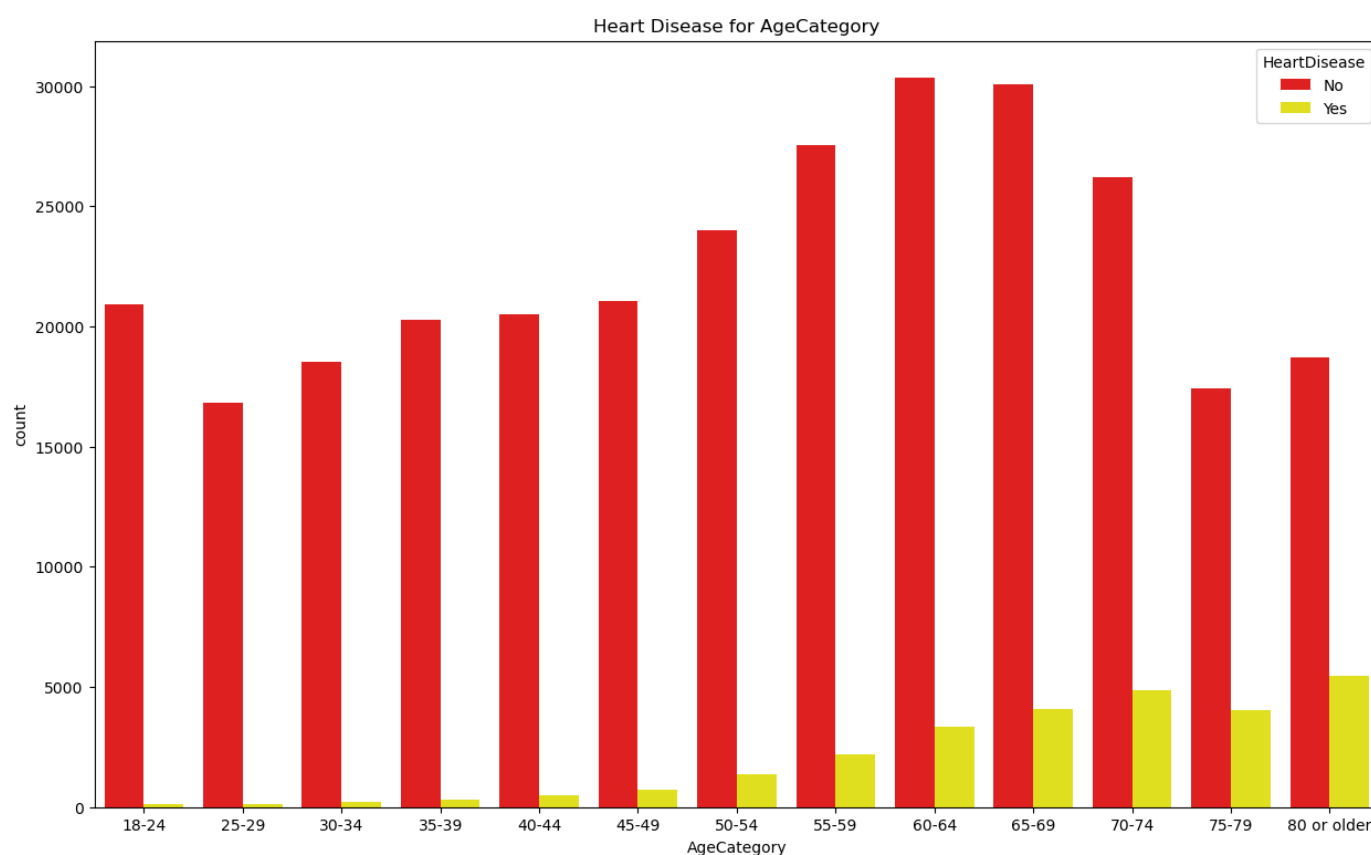
```
In [19]: plt.figure(figsize=(10,5))
sns.countplot(x=new_df['AlcoholDrinking'],data=new_df,hue="HeartDisease",palette=['Red',
plt.title("Heat Disease because of Alcohol Drinking")
```


Out[19]: Text(0.5, 1.0, 'Heart Disease because of Alcohol Drinking')



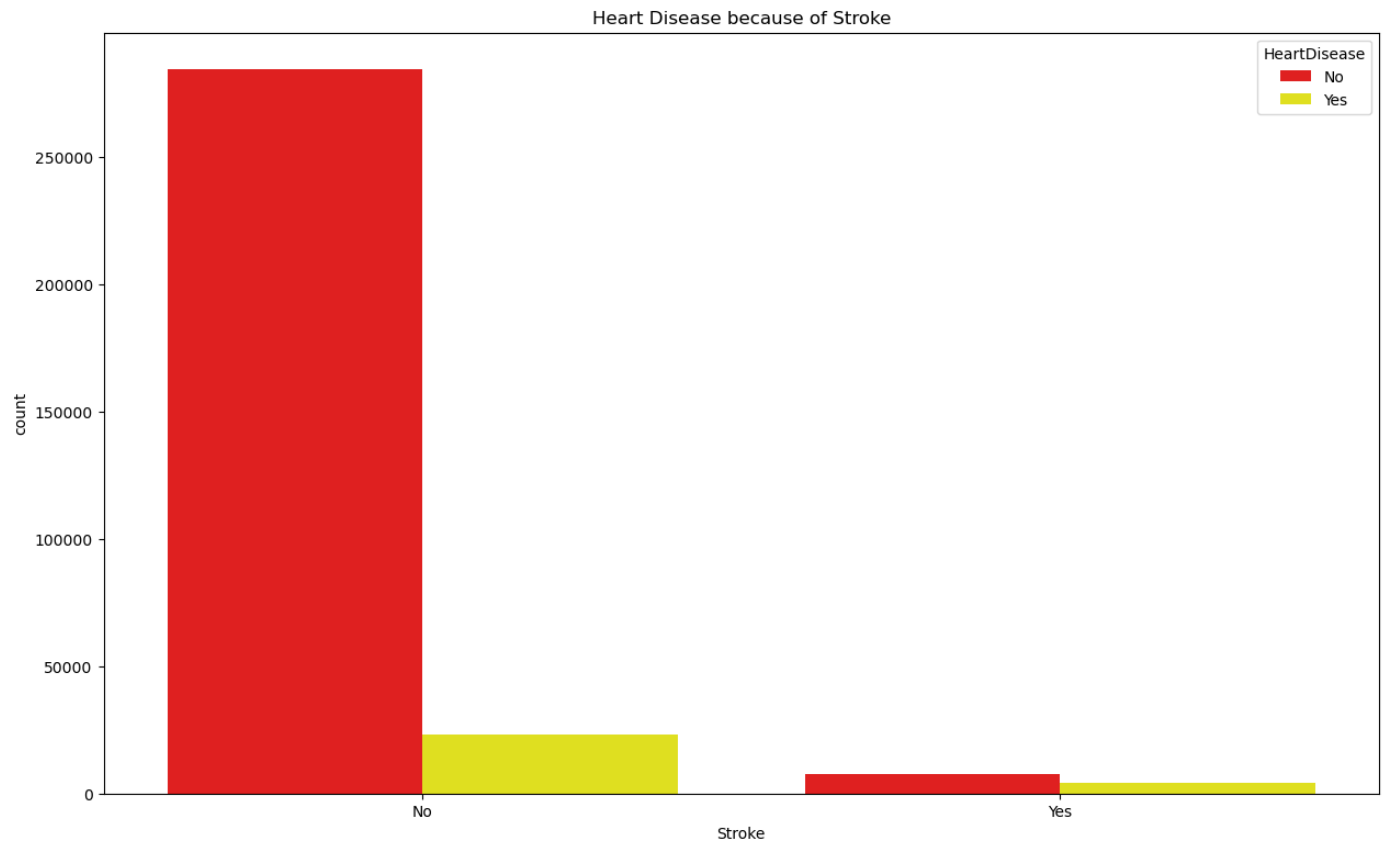
```
In [20]: plt.figure(figsize=(15,9))
sns.countplot(x=new_df['AgeCategory'].sort_values(ascending=True),data=new_df,hue='HeartDisease')
plt.title("Heart Disease for AgeCategory")
```

Out[20]: Text(0.5, 1.0, 'Heart Disease for AgeCategory')



```
In [21]: plt.figure(figsize=(15,9))
sns.countplot(x=new_df['Stroke'].sort_values(ascending=True),data=new_df,hue='HeartDisease')
plt.title("Heart Disease because of Stroke")
```

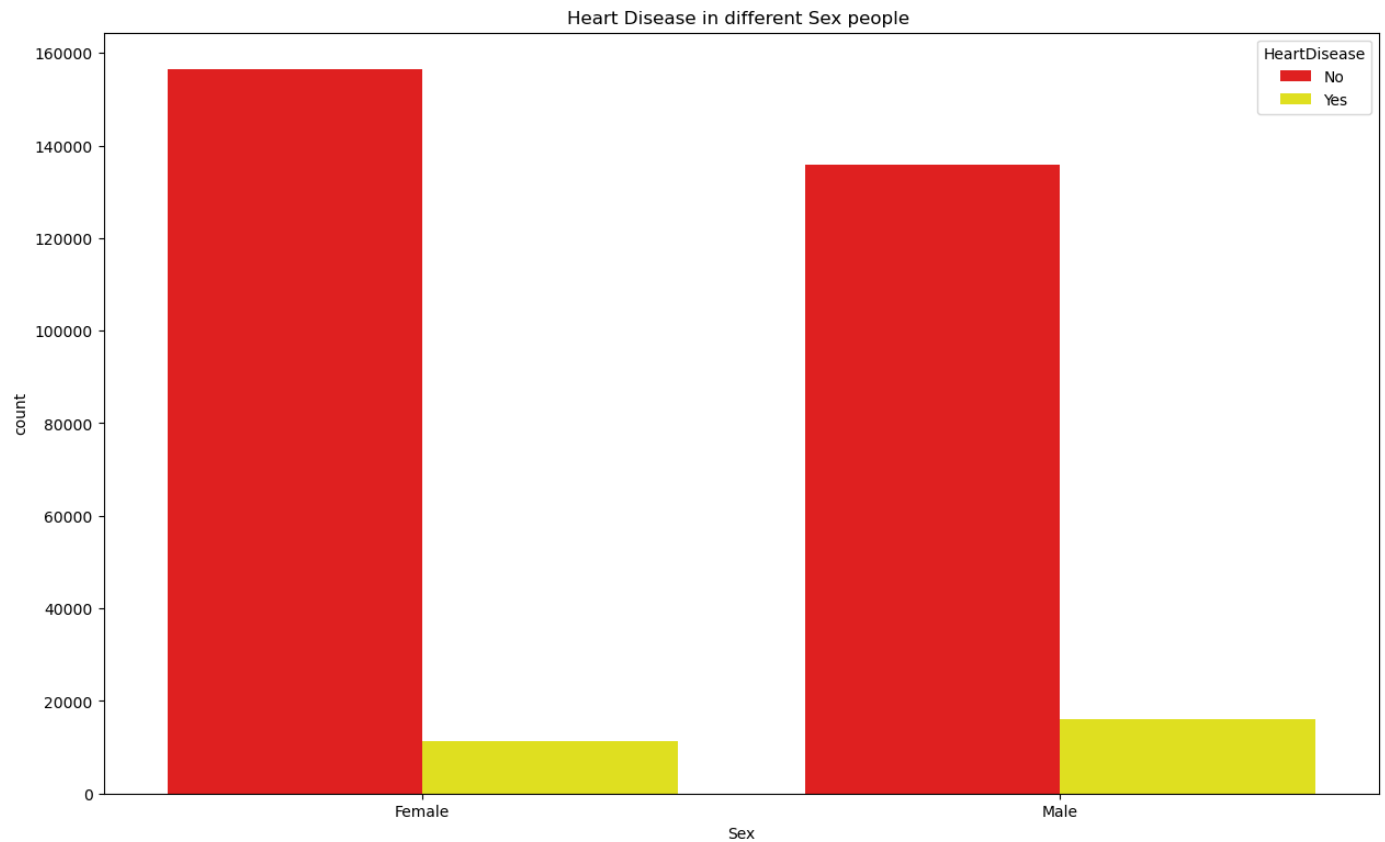
Out[21]: Text(0.5, 1.0, 'Heart Disease because of Stroke')



The frequency of heart disease cases is directly proportional with the age of the individuals, which is expected. Heart disease appears to be more prevalent among individuals who identify as White, indicating the need for more diversity in the racial makeup of the sample group. Insufficient data is available to confirm Alcohol Drinking as a significant indicator of heart disease

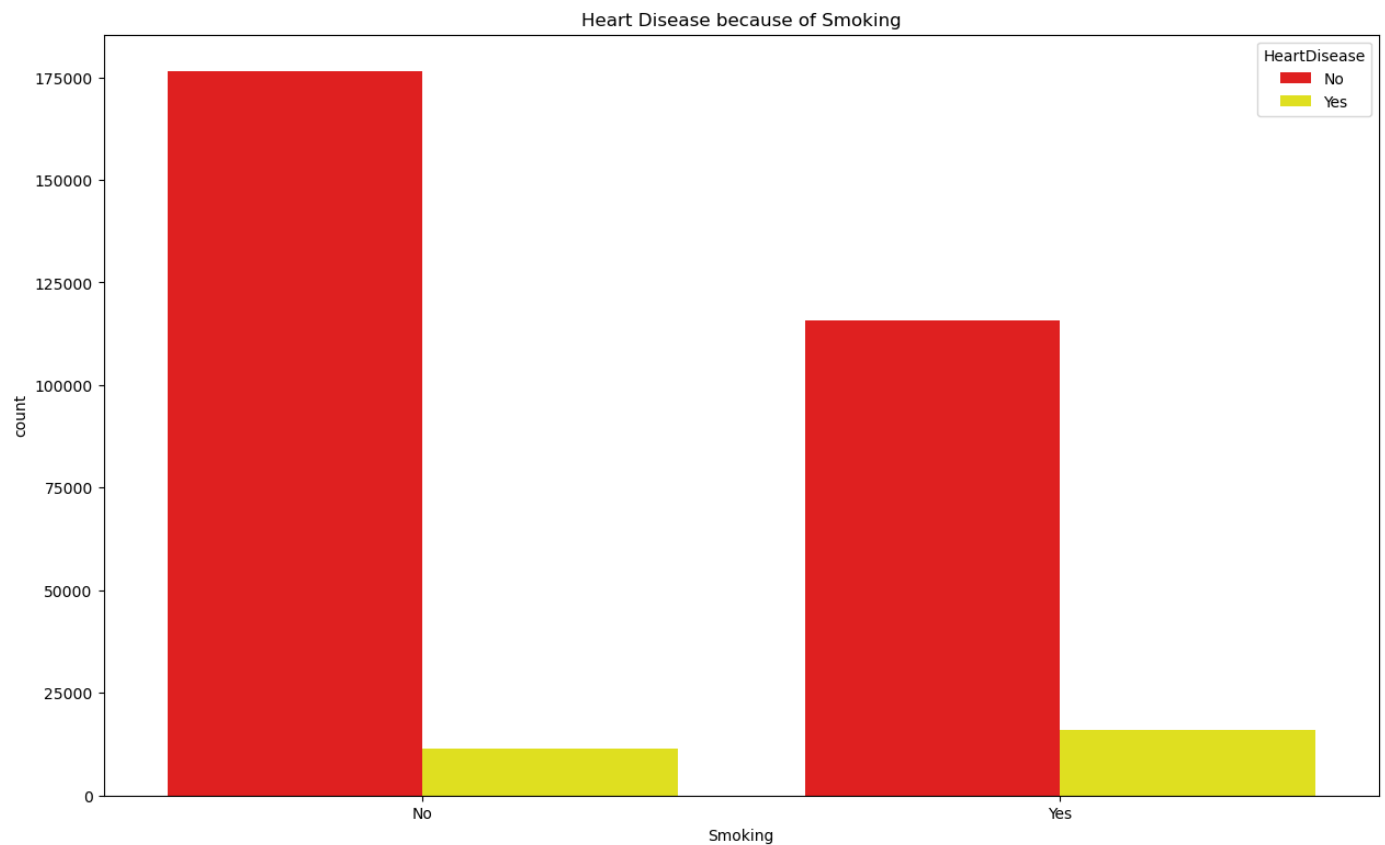
```
In [22]: plt.figure(figsize=(15,9))
sns.countplot(x=new_df['Sex'].sort_values(ascending=True),data=new_df,hue='HeartDisease')
plt.title("Heart Disease in different Sex people ")
```

```
Out[22]: Text(0.5, 1.0, 'Heart Disease in different Sex people ')
```




```
In [23]: plt.figure(figsize=(15,9))
sns.countplot(x=new_df['Smoking'].sort_values(ascending=True),data=new_df,hue='HeartDisease')
plt.title("Heart Disease because of Smoking")
```

```
Out[23]: Text(0.5, 1.0, 'Heart Disease because of Smoking')
```




```
In [24]: sns.countplot(x=new_df['PhysicalActivity'],hue="HeartDisease",data=new_df,palette="winter")
plt.title("Heart Disease due to Physical Acitivity")
```

```
Text(0.5, 1.0, 'Heart Disease due to Physical Acitivity')
```



```
In [25]: sns.countplot(x=new_df['GenHealth'].sort_values(ascending=True),hue="HeartDisease",data=
plt.title("Heart Disease due to General Health")
```

```
Out[25]: Text(0.5, 1.0, 'Heart Disease due to General Health')
```




```
In [26]: sns.distplot(x=new_df['PhysicalHealth'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning:

``distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).`

Out[26]: <AxesSubplot:ylabel='Density'>



```
In [27]: plt.figure(figsize=(10,10))
cor=new_df.corr()
sns.heatmap(cor,annot=True,cmap=plt.cm.Reds,fmt='.3f')
```

Out[27]: <AxesSubplot:>



In []:

as we know bmi<18.5 is underweight
bmi between 18.5-24.9 is normal
bmi between 25-29.9 is overweight
bmi between 30-34.9 is obese
bmi >35 is extremely obese

Segment Bmi

```
In [28]: value,index=new_df['BMI'].value_counts().values,new_df['BMI'].value_counts().index
print("BMI Column's Count Values : ")
pd.DataFrame(value,index,columns=['Count'])
```

BMI Column's Count Values :

Out[28]:	Count
26.63	3762
27.46	2767
27.44	2723
24.41	2696
27.12	2525
...	...
59.85	1
50.59	1
92.53	1
62.95	1
46.56	1

3604 rows × 1 columns

```
In [29]: bins = [0, 18.5, 25, 30, 35, np.inf]
names = ['Underweight', 'Normal weight', 'Overweight', 'Obese', 'Extremly Obese']
new_df['SegmentBMI'] = pd.cut(new_df['BMI'],bins, labels=names)
new_df.drop('BMI',axis=1,inplace=True)
new_df.head()
```

C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2431362878.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2431362878.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[29]:	HeartDisease	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth	DiffWalking	Sex	AgeCate
0	No	Yes	No	No	3.0	30.0	No	Female	5
1	No	No	No	Yes	0.0	0.0	No	Female	80 or c
2	No	Yes	No	No	20.0	30.0	No	Male	6
3	No	No	No	No	0.0	0.0	No	Female	7
4	No	No	No	No	28.0	0.0	Yes	Female	4

```
In [30]: value,index=new_df['SegmentBMI'].value_counts().values,new_df['SegmentBMI'].value_counts
print("SegmentBMI Column's Count Values:")
pd.DataFrame(value,index,columns=['Count'])
```

SegmentBMI Column's Count Values:

Out[30]:

	Count
Overweight	114355
Normal weight	97778
Obese	61169
Extremly Obese	41379
Underweight	5114

```
In [31]: #object Columns
obj=new_df.select_dtypes(include=object).columns
pd.DataFrame(obj,columns=['Object Columns'])
```

Out[31]:

	Object Columns
0	HeartDisease
1	Smoking
2	AlcoholDrinking
3	Stroke
4	DiffWalking
5	Sex
6	AgeCategory
7	Diabetic
8	PhysicalActivity
9	GenHealth
10	Asthma
11	KidneyDisease

```
In [32]: new_df['SegmentBMI'].dtypes
```

Out[32]: CategoricalDtype(categories=['Underweight', 'Normal weight', 'Overweight', 'Obese', 'Extremly Obese'], ordered=True)

```
In [33]: obj=list(obj)
obj.append('SegmentBMI')
pd.DataFrame(obj,columns=['Object Columns'])
```


Out [33]:

Object Columns	
0	HeartDisease
1	Smoking
2	AlcoholDrinking
3	Stroke
4	DiffWalking
5	Sex
6	AgeCategory
7	Diabetic
8	PhysicalActivity
9	GenHealth
10	Asthma
11	KidneyDisease
12	SegmentBMI

```
In [34]: #Transform Object Columns
#label_encoder = preprocessing.LabelEncoder()
label=LabelEncoder()
for col in obj:
    new_df[col]=label.fit_transform(new_df[col])
new_df
```

```
C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\Jkuma\AppData\Local\Temp\ipykernel_1768\2400054535.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[34]:

	HeartDisease	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth	DiffWalking	Sex	AgeCa
0	0	1	0	0	3.0	30.0	0	0	
1	0	0	0	1	0.0	0.0	0	0	
2	0	1	0	0	20.0	30.0	0	1	
3	0	0	0	0	0.0	0.0	0	0	
4	0	0	0	0	28.0	0.0	1	0	
...	
319790	1	1	0	0	7.0	0.0	1	1	
319791	0	1	0	0	0.0	0.0	0	1	
319792	0	0	0	0	0.0	0.0	0	0	
319793	0	0	0	0	0.0	0.0	0	0	
319794	0	0	0	0	0.0	0.0	0	0	

319795 rows × 16 columns

```
In [35]: '''sns.pairplot(data = new_df , hue= 'HeartDisease')
plt.legend('HeartDisease')
plt.show()'''
```

```
Out[35]: "sns.pairplot(data = new_df , hue= 'HeartDisease')\nplt.legend('HeartDisease')\nplt.show()\n"
```

```
In [36]: plt.figure(figsize=(20,10))
sns.heatmap(new_df.corr(), annot=True, cbar=False, cmap='RdBu')
plt.savefig("heatmap.png")
new_df.corr()
```

Out[36]:

	HeartDisease	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth	DiffWalking
HeartDisease	1.000000	0.107764	-0.032080	0.196835	0.170721	0.028591	0.201258
Smoking	0.107764	1.000000	0.111768	0.061226	0.115352	0.085157	0.120074
AlcoholDrinking	-0.032080	0.111768	1.000000	-0.019858	-0.017254	0.051282	-0.035328
Stroke	0.196835	0.061226	-0.019858	1.000000	0.137014	0.046467	0.174143
PhysicalHealth	0.170721	0.115352	-0.017254	0.137014	1.000000	0.287987	0.428373
MentalHealth	0.028591	0.085157	0.051282	0.046467	0.287987	1.000000	0.152235
DiffWalking	0.201258	0.120074	-0.035328	0.174143	0.428373	0.152235	1.000000
Sex	0.070040	0.085052	0.004200	-0.003091	-0.040904	-0.100058	-0.068860
AgeCategory	0.233432	0.128331	-0.059528	0.137822	0.110763	-0.155506	0.243263
Diabetic	0.168553	0.053847	-0.057372	0.101518	0.151361	0.032945	0.205502
PhysicalActivity	-0.100030	-0.097174	0.017487	-0.079455	-0.232283	-0.095808	-0.278524
GenHealth	-0.011062	0.020625	0.001629	-0.009335	-0.035703	-0.004412	-0.043552
SleepTime	0.008327	-0.030336	-0.005065	0.011900	-0.061387	-0.119717	-0.022216
Asthma	0.041444	0.024149	-0.002202	0.038866	0.117907	0.114008	0.103222
KidneyDisease	0.145197	0.034920	-0.028280	0.091167	0.142197	0.037281	0.153064
SegmentBMI	0.002387	0.013652	0.008293	0.002065	-0.051750	-0.059308	-0.078957

HeartDisease	1	0.11	-0.032	0.2	0.17	0.029	0.2	0.07	0.23	0.17	-0.1	-0.011	0.0083	0.041	0.15	0.0024
Smoking	0.11	1	0.11	0.061	0.12	0.085	0.12	0.085	0.13	0.054	-0.097	0.021	-0.03	0.024	0.035	0.014
AlcoholDrinking	-0.032	0.11	1	-0.02	-0.017	0.051	-0.035	0.0042	-0.06	-0.057	0.017	0.0016	-0.0051	-0.0022	-0.028	0.0083
Stroke	0.2	0.061	-0.02	1	0.14	0.046	0.17	-0.0031	0.14	0.1	-0.079	-0.0093	0.012	0.039	0.091	0.0021
PhysicalHealth	0.17	0.12	-0.017	0.14	1	0.29	0.43	-0.041	0.11	0.15	-0.23	-0.036	-0.061	0.12	0.14	-0.052
MentalHealth	0.029	0.085	0.051	0.046	0.29	1	0.15	-0.1	-0.16	0.033	-0.096	-0.0044	-0.12	0.11	0.037	-0.059
DiffWalking	0.2	0.12	-0.035	0.17	0.43	0.15	1	-0.069	0.24	0.21	-0.28	-0.044	-0.022	0.1	0.15	-0.079
Sex	0.07	0.085	0.0042	-0.0031	-0.041	-0.1	-0.069	1	-0.067	-0.013	0.048	-0.01	-0.016	-0.069	-0.0091	0.11
AgeCategory	0.23	0.13	-0.06	0.14	0.11	-0.16	0.24	-0.067	1	0.19	-0.12	0.044	0.1	-0.058	0.12	0.064
Diabetic	0.17	0.054	-0.057	0.1	0.15	0.033	0.21	-0.013	0.19	1	-0.13	-0.011	0.00045	0.05	0.14	-0.062
PhysicalActivity	-0.1	-0.097	0.017	-0.079	-0.23	-0.096	-0.28	0.048	-0.12	-0.13	1	0.024	0.0038	-0.042	-0.082	0.059
GenHealth	-0.011	0.021	0.0016	-0.0093	-0.036	-0.0044	-0.044	-0.01	0.044	-0.011	0.024	1	-0.0042	0.0073	-0.011	0.034
SleepTime	0.0083	-0.03	-0.0051	0.012	-0.061	-0.12	-0.022	-0.016	0.1	0.00045	0.0038	-0.0042	1	-0.048	0.0062	0.013
Asthma	0.041	0.024	-0.0022	0.039	0.12	0.11	0.1	-0.069	-0.058	0.05	-0.042	0.0073	-0.048	1	0.04	-0.05
KidneyDisease	0.15	0.035	-0.028	0.091	0.14	0.037	0.15	-0.0091	0.12	0.14	-0.082	-0.011	0.0062	0.04	1	-0.016
SegmentBMI	0.0024	0.014	0.0083	0.0021	-0.052	-0.059	-0.079	0.11	0.064	-0.062	0.059	0.034	0.013	-0.05	-0.016	1

```
In [37]: X=new_df.drop(['HeartDisease'],axis=1)
y=new_df['HeartDisease']
```

```
In [38]: X.head()
```

```
Out[38]:
```

	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth	DiffWalking	Sex	AgeCategory	Diabetic	PhysicalActivity
0	1	0	0	3.0	30.0	0	0	7	2	
1	0	0	1	0.0	0.0	0	0	12	0	
2	1	0	0	20.0	30.0	0	1	9	2	
3	0	0	0	0.0	0.0	0	0	11	0	
4	0	0	0	28.0	0.0	1	0	4	0	

```
In [39]: y.head()
```

```
Out[39]:
```

0	0
1	0
2	0
3	0
4	0

Name: HeartDisease, dtype: int32

```
In [40]: from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import BaggingClassifier, RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, f1_score,
```

NearMiss(ubdersampling)

```
In [42]: nm=NearMiss()
X_res,y_res=nm.fit_resample(X,y)
X_res.shape,y_res.shape
```

```
Out[42]: ((54746, 15), (54746,))
```

```
In [43]: Xtrain,Xtest,ytrain,ytest = train_test_split(X_res, y_res, test_size = 0.2, random_state
```

```
In [44]: bag_clf = BaggingClassifier(LogisticRegression(), n_estimators = 100, bootstrap = True,n
bag_clf.fit(Xtrain, ytrain)
bag_clf.oob_score_
```


```
Out[44]: 0.8776372271440315
```

```
In [45]: y_pred_lr = bag_clf.predict(Xtest)

# Create the confusion matrix
cm = confusion_matrix(ytest, y_pred_lr)
#print(confusion_matrix_rf)

TN, FP, FN, TP = cm.ravel()
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'], in
sns.heatmap(cm_matrix, annot=True, fmt='d')
plt.title('Confusion Matrix for logistic Regression', fontsize=14, fontweight='bold')
```

```
Out[45]: Text(0.5, 1.0, 'Confusion Matrix for logistic Regression')
```



```
In [46]: print("Classification Report of Test Dataset Logistics Regression\n")
print(classification_report(ytest, y_pred_lr))
```

Classification Report of Test Dataset Logistics Regression


	precision	recall	f1-score	support
0	0.84	0.94	0.88	5441
1	0.93	0.82	0.87	5509
accuracy			0.88	10950
macro avg	0.88	0.88	0.88	10950
weighted avg	0.88	0.88	0.88	10950

```
In [47]: y_prob_test = bag_clf.predict_proba(Xtest)[: ,1]
fpr, tpr, thresholds = roc_curve(ytest, y_pred_lr)

auc = round(roc_auc_score(ytest, y_pred_lr), 2)

plt.plot(fpr, tpr, linewidth=2, label="ROC curve (auc=" + str(auc) + ")")
plt.plot([0,1], [0,1], 'k--' )
plt.title('ROC curve for Predicting Heart Disease of Logistic Regression', fontsize=14,
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.show()
```

ROC curve for Predicting Heart Disease of Logistic Regression



```
In [48]: # Evaluate model on training dataset
train_accuracy = accuracy_score(ytrain, bag_clf.predict(Xtrain))
print("Training Accuracy for logistic regression:", round(train_accuracy, 2))

# Evaluate model on testing dataset
test_accuracy = accuracy_score(ytest, bag_clf.predict(Xtest))
print("Testing Accuracy for logistic regression:", round(test_accuracy, 2))
```

Training Accuracy for logistic regression: 0.88

Testing Accuracy for logistic regression: 0.88

```
In [49]: ypred = bag_clf.predict(Xtest)
print("Accuracy Score : ", accuracy_score(ytest, ypred))
print("Precision Score : ", precision_score(ytest, ypred))
print("Recall Score : ", recall_score(ytest, ypred))
print("F1 Score : ", f1_score(ytest, ypred))
```

Accuracy Score : 0.8783561643835617
Precision Score : 0.9333886698485163
Recall Score : 0.8164821201669995
F1 Score : 0.8710302091402015

```
In [50]: bag_clf = BaggingClassifier(DecisionTreeClassifier(), n_estimators = 50,  
                                     bootstrap = True,  
                                     n_jobs = -1, oob_score = True)  
bag_clf.fit(Xtrain, ytrain)  
bag_clf.oob_score_
```


```
Out[50]: 0.9212256827107499
```

```
In [51]: y_pred_df = bag_clf.predict(Xtest)  
  
# Create the confusion matrix  
confusion_matrix_rf = confusion_matrix(ytest, y_pred_df)  
print(confusion_matrix_rf)
```

```
[[5360  81]  
 [ 748 4761]]
```

```
In [52]: y_pred_df = bag_clf.predict(Xtest)  
  
# Create the confusion matrix  
cm = confusion_matrix(ytest, y_pred_df)  
#print(confusion_matrix_rf)  
  
TN, FP, FN, TP = cm.ravel()  
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'], in  
sns.heatmap(cm_matrix, annot=True, fmt='d')  
plt.title('Confusion Matrix for Decision Tree', fontsize=14, fontweight='bold')
```

```
Out[52]: Text(0.5, 1.0, 'Confusion Matrix for Decision Tree')
```



```
In [53]: print("Classification Report of Test Dataset for Decision Tree\n")  
print(classification_report(ytest, y_pred_df))
```



	precision	recall	f1-score	support
0	0.88	0.99	0.93	5441
1	0.98	0.86	0.92	5509
accuracy			0.92	10950
macro avg	0.93	0.92	0.92	10950
weighted avg	0.93	0.92	0.92	10950

```
In [54]: y_prob_test = bag_clf.predict_proba(Xtest)[: ,1]
fpr, tpr, thresholds = roc_curve(ytest, y_pred_df)

auc = round(roc_auc_score(ytest, y_pred_df), 2)

plt.plot(fpr, tpr, linewidth=2, label="ROC curve (auc=" + str(auc) + ")")
plt.plot([0,1], [0,1], 'k--' )
plt.title('ROC curve for Predicting Heart Disease of decision tree', fontsize=14, fontwe
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.show()
```

ROC curve for Predicting Heart Disease of decision tree



```
In [55]: # Evaluate model on training dataset
train_accuracy = accuracy_score(ytrain, bag_clf.predict(Xtrain))
print("Training Accuracyfor Decision Tree:", round(train_accuracy, 2))

# Evaluate model on testing dataset
test_accuracy = accuracy_score(ytest, bag_clf.predict(Xtest))
print("Testing Accuracyfor Decision Tree:", round(test_accuracy, 2))
```

Training Accuracyfor Decision Tree: 0.93

Testing Accuracyfor Decision Tree: 0.92

```
In [56]: ypred = bag_clf.predict(Xtest)
print("Accuracy Score : ", accuracy_score(ytest, ypred))
print("Precision Score : ", precision_score(ytest, ypred))
print("Recall Score : ", recall_score(ytest, ypred))
print("F1 Score : ", f1_score(ytest, ypred))
```

```
Accuracy Score : 0.9242922374429223
Precision Score : 0.983271375464684
Recall Score : 0.8642221818841895
F1 Score : 0.9199111196985799
```

```
In [57]: rnd_clf = RandomForestClassifier(n_estimators = 50, n_jobs = -1, oob_score= True)
rnd_clf.fit(Xtrain, ytrain)
rnd_clf.oob_score_
```

```
Out[57]: 0.9217965110969039
```

```
In [58]: y_pred_rf = bag_clf.predict(Xtest)

# Create the confusion matrix
confusion_matrix_rf = confusion_matrix(ytest, y_pred_rf)
print(confusion_matrix_rf)
```


```
[[5360  81]
 [ 748 4761]]
```

```
In [59]: y_pred_rf = bag_clf.predict(Xtest)

# Create the confusion matrix
cm = confusion_matrix(ytest, y_pred_rf)
#print(confusion_matrix_rf)

TN, FP, FN, TP = cm.ravel()
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'], in
sns.heatmap(cm_matrix, annot=True, fmt='d')
plt.title('Confusion Matrix for Random Forest', fontsize=14, fontweight='bold')
```

```
Out[59]: Text(0.5, 1.0, 'Confusion Matrix for Random Forest')
```



```
In [60]: print("Classification Report of Test Dataset for Random Forest\n")
print(classification_report(ytest, y_pred_rf))
```

Classification Report of Test Dataset for Random Forest


	precision	recall	f1-score	support
0	0.88	0.99	0.93	5441
1	0.98	0.86	0.92	5509
accuracy			0.92	10950
macro avg	0.93	0.92	0.92	10950
weighted avg	0.93	0.92	0.92	10950

```
In [61]: y_prob_test = bag_clf.predict_proba(Xtest)[:,-1]
fpr, tpr, thresholds = roc_curve(ytest, y_pred_rf)

auc = round(roc_auc_score(ytest, y_pred_rf), 2)

plt.plot(fpr, tpr, linewidth=2, label="ROC curve (auc=" + str(auc) + ")")
plt.plot([0,1], [0,1], 'k--')
plt.title('ROC curve for Predicting Heart Disease using Random Forest', fontsize=14, font
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.show()
```

ROC curve for Predicting Heart Disease using Random Forest



```
In [62]: ypred = rnd_clf.predict(Xtest)
print("Accuracy Score : ", accuracy_score(ytest, ypred))
print("Precision Score : ", precision_score(ytest, ypred))
print("Recall Score : ", recall_score(ytest, ypred))
print("F1 Score : ", f1_score(ytest, ypred))
```

```
Accuracy Score : 0.9252054794520548
Precision Score : 0.9841040462427746
Recall Score : 0.8653113087674714
F1 Score : 0.920892494929006
```

```
In [63]: # Evaluate model on training dataset
train_accuracy = accuracy_score(ytrain, rnd_clf.predict(Xtrain))
print("Training Accuracy for Random Forest:", round(train_accuracy, 2))

# Evaluate model on testing dataset
test_accuracy = accuracy_score(ytest, rnd_clf.predict(Xtest))
print("Testing Accuracy for Random Forest:", round(test_accuracy, 2))
```

```
Training Accuracy for Random Forest: 0.93
Testing Accuracy for Random Forest: 0.93
```

SMOTE (oversampling)

```
In [64]: from imblearn.over_sampling import SMOTE
sm = SMOTE(random_state = 0)
sm.fit(X,y)
x_resem, y_resem = sm.fit_resample(X, y)
```

```
In [65]: x_resem.shape, y_resem.shape
```

```
Out[65]: ((584844, 15), (584844,))
```

```
In [66]: y_resem.value_counts()
```

```
Out[66]: 0    292422  
         1    292422  
         Name: HeartDisease, dtype: int64
```

```
In [67]: Xtrain,Xtest,ytrain,ytest = train_test_split(x_resem, y_resem, test_size = 0.2, random_s
```


```
In [68]: bag_clf = BaggingClassifier(LogisticRegression(), n_estimators = 100, bootstrap = True,n  
bag_clf.fit(Xtrain, ytrain)  
bag_clf.oob_score_
```

```
Out[68]: 0.7344461661768634
```

```
In [69]: y_pred_lr = bag_clf.predict(Xtest)  
  
# Create the confusion matrix  
cm = confusion_matrix(ytest, y_pred_lr)  
#print(confusion_matrix_rf)  
  
TN, FP, FN, TP = cm.ravel()  
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'], in  
sns.heatmap(cm_matrix, annot=True, fmt='d')  
plt.title('Confusion Matrix for logistic Regression', fontsize=14, fontweight='bold')
```

```
Out[69]: Text(0.5, 1.0, 'Confusion Matrix for logistic Regression')
```

Confusion Matrix for logistic Regression



```
In [70]: print("Classification Report of Test Dataset for logistic Regression\n")  
print(classification_report(ytest, y_pred_lr))
```


	precision	recall	f1-score	support
0	0.75	0.70	0.72	58461
1	0.72	0.77	0.74	58508
accuracy			0.73	116969
macro avg	0.73	0.73	0.73	116969
weighted avg	0.73	0.73	0.73	116969

```
In [71]: y_prob_test = bag_clf.predict_proba(Xtest)[:,-1]
fpr, tpr, thresholds = roc_curve(ytest, y_pred_lr)

auc = round(roc_auc_score(ytest, y_pred_lr), 2)

plt.plot(fpr, tpr, linewidth=2, label="ROC curve (auc=" + str(auc) + ")")
plt.plot([0,1], [0,1], 'k--' )
plt.title('ROC curve for Predicting Heart Disease for logistic Regression', fontsize=14,
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.show()
```

ROC curve for Predicting Heart Disease for logistic Regression



```
In [72]: ypred = bag_clf.predict(Xtest)
print("Accuracy Score : ", accuracy_score(ytest, ypred))
print("Precision Score : ", precision_score(ytest, ypred))
print("Recall Score : ", recall_score(ytest, ypred))
print("F1 Score : ", f1_score(ytest, ypred))
```

```
Accuracy Score : 0.732766801460216
Precision Score : 0.7180209300092809
Recall Score : 0.7669378546523552
F1 Score : 0.7416736913439448
```

```
In [73]: # Evaluate model on training dataset
train_accuracy = accuracy_score(ytrain, bag_clf.predict(Xtrain))
print("Training Accuracy for logistic Regression:", round(train_accuracy, 2))

# Evaluate model on testing dataset
test_accuracy = accuracy_score(ytest, bag_clf.predict(Xtest))
print("Testing Accuracy for logistic Regression:", round(test_accuracy, 2))
```

Training Accuracy for logistic Regression: 0.73
Testing Accuracy for logistic Regression: 0.73

```
In [74]: bag_clf = BaggingClassifier(DecisionTreeClassifier(), n_estimators = 50,
                                   bootstrap = True,
                                   n_jobs = -1, oob_score = True)

bag_clf.fit(Xtrain, ytrain)
bag_clf.oob_score_
```


Out[74]: 0.8678706919583222

```
In [75]: y_pred_df = bag_clf.predict(Xtest)

# Create the confusion matrix
cm = confusion_matrix(ytest, y_pred_df)
#print(confusion_matrix_rf)

TN, FP, FN, TP = cm.ravel()
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'], in
sns.heatmap(cm_matrix, annot=True, fmt='d')
plt.title('Confusion Matrix for Decision Tree', fontsize=14, fontweight='bold')
```

Out[75]: Text(0.5, 1.0, 'Confusion Matrix for Decision Tree')



```
In [76]: print("Classification Report of Test Dataset for Decision Tree\n")
print(classification_report(ytest, y_pred_df))
```

	precision	recall	f1-score	support
0	0.89	0.84	0.86	58461
1	0.85	0.90	0.87	58508
accuracy			0.87	116969
macro avg	0.87	0.87	0.87	116969
weighted avg	0.87	0.87	0.87	116969

```
In [77]: train_accuracy = accuracy_score(ytrain, bag_clf.predict(Xtrain))
print("Training Accuracy for Random Forest:", round(train_accuracy, 2))
```

```
# Evaluate model on testing dataset
test_accuracy = accuracy_score(ytest, bag_clf.predict(Xtest))
print("Testing Accuracy for Random Forest:", round(test_accuracy, 2))
```


Training Accuracy for Random Forest: 0.92
Testing Accuracy for Random Forest: 0.87

```
In [78]: y_prob_test = bag_clf.predict_proba(Xtest)[:,-1]
fpr, tpr, thresholds = roc_curve(ytest, y_pred_df)

auc = round(roc_auc_score(ytest, y_pred_df), 2)

plt.plot(fpr, tpr, linewidth=2, label="ROC curve (auc=" + str(auc) + ")")
plt.plot([0,1], [0,1], 'k--')
plt.title('ROC curve for Predicting Heart Disease of decision tree', fontsize=14, fontwe
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.show()
```

ROC curve for Predicting Heart Disease of decision tree




```
In [79]: ypred = bag_clf.predict(Xtest)
print("Accuracy Score : ", accuracy_score(ytest, ypred))
print("Precision Score : ", precision_score(ytest, ypred))
print("Recall Score : ", recall_score(ytest, ypred))
print("F1 Score : ", f1_score(ytest, ypred))
```

```
Accuracy Score : 0.8675717497798562
Precision Score : 0.8458483406225881
Recall Score : 0.899107814315991
F1 Score : 0.8716652858326429
```

```
In [80]: rnd_clf = RandomForestClassifier(n_estimators = 50, n_jobs = -1, oob_score= True)
rnd_clf.fit(Xtrain, ytrain)
rnd_clf.oob_score_
```


```
Out[80]: 0.8680138925995191
```

```
In [81]: y_pred_rf = bag_clf.predict(Xtest)

# Create the confusion matrix
cm = confusion_matrix(ytest, y_pred_rf)
#print(confusion_matrix_rf)

TN, FP, FN, TP = cm.ravel()
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'], in
sns.heatmap(cm_matrix, annot=True, fmt='d')
plt.title('Confusion Matrix for Random Forest', fontsize=14, fontweight='bold')
```

```
Out[81]: Text(0.5, 1.0, 'Confusion Matrix for Random Forest')
```



```
In [82]: print("Classification Report of Test Dataset for Random Forest\n")
print(classification_report(ytest, y_pred_rf))
```


	precision	recall	f1-score	support
0	0.89	0.84	0.86	58461
1	0.85	0.90	0.87	58508
accuracy			0.87	116969
macro avg	0.87	0.87	0.87	116969
weighted avg	0.87	0.87	0.87	116969

```
In [83]: y_prob_test = bag_clf.predict_proba(Xtest)[:,-1]
fpr, tpr, thresholds = roc_curve(ytest, y_pred_rf)

auc = round(roc_auc_score(ytest, y_pred_rf), 2)

plt.plot(fpr, tpr, linewidth=2, label="ROC curve (auc=" + str(auc) + ")")
plt.plot([0,1], [0,1], 'k--')
plt.title('ROC curve for Predicting Heart Disease using Random Forest', fontsize=14, fontweight='bold')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.show()
```

ROC curve for Predicting Heart Disease using Random Forest



```
In [84]: ypred = rnd_clf.predict(Xtest)
print("Accuracy Score : ", accuracy_score(ytest, ypred))
print("Precision Score : ", precision_score(ytest, ypred))
print("Recall Score : ", recall_score(ytest, ypred))
print("F1 Score : ", f1_score(ytest, ypred))
```

```
Accuracy Score : 0.8690165770417803
Precision Score : 0.8462994146419693
Recall Score : 0.9019450331578588
F1 Score : 0.8732366398318758
```

```
In [85]: # Evaluate model on training dataset
train_accuracy = accuracy_score(ytrain, bag_clf.predict(Xtrain))
print("Training Accuracy for Random Forest:", round(train_accuracy, 2))

# Evaluate model on testing dataset
test_accuracy = accuracy_score(ytest, bag_clf.predict(Xtest))
print("Testing Accuracy for Random Forest:", round(test_accuracy, 2))
```

Training Accuracy for Random Forest: 0.92
Testing Accuracy for Random Forest: 0.87

In []:

In []: